Comparison of TCP and SCTP transport protocols over wireless multi-hop networks

Abstract

The Stream Control Transmission Protocol (SCTP) was defined, about twenty years after TCP and UDP, in 2000 by the Internet Engineering Task Force (IETF). They combined the best practices from both older protocols to create a message-based, multi-streamed transport protocol. There are papers [1][2] which shows the improvements of the SCTP with the usage of HTTP over the internet. In this paper the behaviour of the TCP and SCTP in a wireless multi-hop network environment is compared. For this the OMNeT++ network simulation framework [3] has been used. The hops are in motion (different speeds) and use the proactive OLSR protocol to transmit data over a wide open field (500m x 500m) without obstacles like mountains or buildings. In the first part the reader is introduced into some technical topics for a better understanding of the used technologies. The simulation setting and the corresponding results are part of the mid section. At the end of the document a conclusion of the project with a summary of the observations is made.

Contents

1   Introduction

1.1   Transmission Control Protocol

The Transmission Control Protocol (TCP) is besides the User Datagram Protocol (UDP) one of the core protocols of the Internet Protocol suite (IP) which has been developed between the mid-seventies and the early-eighties of the 20th century.
TCP provides connection oriented data-transfer with the following key features:

- Ordered data transfer

- Retransmission of lost packets

- Error-free data transfer

- Flow control

- Congestion control

Since TCP is the most common used transport protocol it is assumed that the reader is familiar with the operation of the Protocol. Because of that the focus of the introduction lies on the less common protocols used in this simulation scenario.

1.2   Stream Control Transmission Protocol

Stream Control Transmission Protocol (SCTP) is a reliable, message oriented transport protocol that provides new services and features for IP communication. For the past twenty years, reliable communication service has been provided by TCP, unreliable by UDP. So, what has brought about the addition of a third protocol to the IP suite of protocols? Many of the features found in TCP and UDP can also be found in SCTP (see Table 1).

Service/Features	SCTP	TCP	UDP
Message-Oriented	yes	no	yes
Byte-Oriented	no	yes	no
Connection-Oriented	yes	yes	no
Full Duplex	yes	yes	yes
Reliable data transfer	yes	yes	no
Partially-Reliable data transfer	opt	no	no
Ordered data delivery	yes	yes	no
Unordered delivery	yes	no	yes
Flow control	yes	yes	no
Congestion Control	yes	yes	no
Selective Acknowledgements	yes	opt	no
Multistreaming	yes	no	no
Multihoming	yes	no	no
Dynamic Multihoming	opt	no	no
SYN flooding attack prevention	yes	no	n/a
Allows half-closed state	no	yes	n/a
Reach-ability check	yes	opt	no

Table 1: Feature comparison

1.2.1   Key Benefits

SCTP improves upon TCP and UDP by integrating components of each. But the designers of SCTP did not stop there. There have been two new concepts added: multi-homing and multi-streaming.

Multi-homing   SCTP was designed to handle the signalling of telecommunications over IP. Since telecommunications are very susceptible to time delays, every millisecond counts. Multi-homing enables systems that have multiple interfaces, for redundancy, to use one over the other without having to wait. Within SCTP one interface is established as the primary and the rest become secondary. If the primary should fail for whatever reason, a secondary is selected and utilized. When the primary becomes available again, the communications can be transferred back without the application being aware there was an issue. While establishing the connections, the primary and secondary interfaces are checked and monitored using a heartbeat/heartbeat acknowledgement process that validates addresses, and maintains a Round Trip Time (RTT) calculation for each address. The RTT can indicate that the primary is slower than a secondary and allow for the communications to migrate to the secondary interface.

Multi-streaming   Using TCP, only one single data stream is allowed per connection. All of the information must be passed through that one stream. SCTP allows multiple simultaneous data streams within a connection or association. Each message sent to a data stream can have a different final destination, but each must maintain message boundaries. For example, systems cannot send parts of the same message through different streams; one message must go through one stream. When running an ordered data delivery system, if one of the packets is out of order or missing, the stream is blocked pending resolution to the order. This is called "Head-of-Line Blocking." With the use of multi-streams, only the stream that is affected would be blocked; the other streams would continue to flow. By using multi-streaming with SCTP, the issue with web browsers only having the ability to handle two simultaneous connections goes away. The client or the web server could immediately open additional streams and send pictures, text, etc. through each stream, reducing overall latency. This could also reduce overhead that servers often incur with the numerous separate connections required to fulfil a request.

Selective acknowledgements   In standard TCP, every message, or packet of information must be accounted for, resent as necessary, and processed in the order they were sent. SCTP has the ability to selectively acknowledge receipt of missing, disordered, or duplicated messages. Due to the nature of telecommunications most applications would end up discarding any unsynchronized messages. Therefore, the need to send and receive the information is forgone. This would mean that a portion of a word, a portion of a video, or a piece of the whiteboard refresh would be skipped over. The applications and users may notice a slight skip in the voice, video, or refresh. This is referred to as jitter within the telecommunications world and a small amount of jitter is often preferred to having the packet resent and reprocessed which would double the amount of jitter, usually making it more noticeable to the users.

Unordered data delivery   Due to the very nature of networks not all packets may travel across the exact same path. If there is a time-delay using one path over another, the original messages could be out of order when received. Unordered data delivery allows for this instance and can correct the issue by reordering the messages correctly. Using TCP's reliable data transfer feature requires that packets be processed in order. If one is missing or out of order, the packet must be reordered before processing can continue. SCTP allows for unordered data delivery and since it has multiple streams, only the one affected is temporarily blocked.

1.3   Optimized Link State Protocol

Optimized Link State Protocol (OLSR) is a proactive routing protocol, so the routes are always immediately available when needed. OLSR is an optimization version of a pure link state protocol. The topological changes cause the flooding of the topological information to all available hosts in the network. To reduce the possible overhead in the network protocol uses Multipoint Relays (MPR). The idea of MPR is to reduce flooding of broadcasts by reducing the same broadcast in some regions in the network.

1.3.1   Control messages

OLSR uses two kinds of the control messages: Hello and Topology Control (TC). Hello messages are used for finding the information about the link status and the host's neighbours. With the Hello message the Multipoint

Relay (MPR) Selector set is constructed which describes which neighbours has chosen this host to act as MPR and from this information the host can calculate its own set of the MPRs. The Hello messages are sent only one hop away but the TC messages are broadcasted throughout the entire network. TC messages are used for broadcasting information about own advertised neighbours which includes at least the MPR Selector list. The TC messages are broadcasted periodically and only the MPR hosts can forward the TC messages. There is also Multiple Interface Declaration (MID) messages which are used for informing other host that the announcing host can have multiple OLSR interface addresses. The MID message is broadcasted throughout the entire network only by MPRs. There is also a "Host and Network Association" (HNA) message which provides the external routing information by giving the possibility for routing to the external addresses. The HNA message provides information about the network- and the netmask addresses, so that OLSR host can consider that the announcing host can act as a gateway to the announcing set of addresses. The HNA is considered as a generalized version of the TC message with only difference that the TC message can inform about route cancelling while HNA message information is removed only after expiration time.

1.3.2 Multipoint relays

The Multipoint Relays (MPR) is the key idea behind the OLSR protocol to reduce the information exchange overhead. Instead of pure flooding the OLSR uses MPR to reduce the number of the host which broadcasts the information throughout the network. The MPR is a host's one hop neighbour which may forward its messages. The MPR set of host is kept small in order for the protocol to be efficient. In OLSR only the MPRs can forward the data throughout the network. Each host must have the information about the symmetric one hop and two hop neighbours in order to calculate the optimal MPR set. The two hop neighbours are found from the Hello message because each Hello message contains all the hosts' neighbours. Selecting the minimum number of the one hop neighbours which covers all the two hop neighbours is the goal of the MPR selection algorithm. Also each host has the Multipoint Relay Selector set, which indicates which hosts has selected the current host to act as a MPR. When the host gets a new broadcast message, which is need to be spread throughout the network and the message's sender interface address is in the MPR Selector set, then the host must forward the message. Due to the possible changes in the ad hoc network, the MPR Selectors sets are updated continuously using Hello messages.

1.3.3 Advantage

OLSR is also a flat routing protocol, it does not need central administrative system to handle its routing process. The proactive characteristic of the protocol provides that the protocol has all the routing information to all participated hosts in the network. However, as a drawback OLSR protocol needs that each host periodic sends the updated topology information throughout the entire network, this increase the protocols bandwidth usage. But the flooding is minimised by the MPRs, which are only allowed to forward the topological messages.

2 Simulation

2.1 Objectives

OMNeT++ simulation based comparison of the TCP and SCTP transport protocols with respect to performance in wireless multihop networks. Modelling of a point-to-point connection with permanent changing routes due to various moving ad hoc network devices. Simulating an open battlefield area where only free-space loss influences the wireless propagation. Focus the influence of changing routes due to mobility on the announced transport protocols considering various metrics.

2.2 Scenario

An important message including maps and pictures has to be transmitted from the army command center (server) at the outer left center of the battlefield to a special force unit (client) at the outer right center of the battlefield (see Figure 1). Due to war actions the wired connections are broken and the satellite base station has been destroyed by the enemy. The armed forces only have the possibilities of building an ad hoc wireless network between different squads of their troops. They build network based on the moving velocity of the available squads. There are various troop genres with different average velocities and rate of turn:

- Infantry (5 km/h), fast turning

- Jeeps (5 - 30 km/h), fast turning

- Tanks (5 - 30 km/h), medium turning

Each squad / vehicle has only one wireless device which acts as network routing hop. Tested networks will only be built by equal troop genres to avoid mixed moving parameters.

2.3 Experiment design

2.3.1 Topology

In this subsection the technical view of the used topology is introduced. Information proposed in the scenario are specified in a more proper form. The following graphic shows the topology used in the simulation. The network consists of two stationary entities with client and server roles. These stationary devices are interconnected over 10 to 15 moving wireless ad hoc devices building a wireless multihop link in between.


Figure 1: Network Topology

The following list collects the exact specifications:

- Area: Square, 500m x 500m, Free-Space loss

- Number of moving devices: 10

- Number of stationary devices: 2

- Moving speed of mobile devices: 5, 10, 20, 30, 70 km/h

- Energy management of mobile devices: Will not be considered

- Movement: Random Walk, change of direction with different turning speeds

- IEEE802.11a wireless ad hoc network

  - Carrier frequency: 2.4 GHz band

  - Base-Bitrate: 6Mbps

  - Max-Bitrate: 54Mbps

  - Transmitter power: 4.5mW

  - Path loss coefficient: 2 (free space)

- Ad hoc routing protocol: OLSR

- Traffic modelling:

  - TCP: Single session data transfer of 400MiB

  - SCTP: Single session data transfer with 1-5 associated streams

The reasons to use the OLSR protocol in the simulation scenario are:

- Army squads move usally in formation. The distances between the mobile devices and the distribution of them will be constant. The one and two-hop routing table will change slightly.

- Immediate transmission of data. The commander in the army command center will have the new observations as fast as possible. OLSR is ready after a few seconds used for setting up the routing tables. Then everything is done for a immediate transmission.

- Shortest path. Establishment of routing entries for each node in the network using Dijkstra's Shortest Path algorithm.

2.3.2 Simulation

The simulation runtime was 200 seconds. The simulation time limits the maximal covered walk of a mobile node which is necessary because of the limited area.

2.3.3 Metrics

For the comparison of the two transport protocols the following metrics have been used:

- Throughput: The average number of bits delivered over a communication channel within a certain amount of time. This metric has the unit bits per second (bps)

- Round trip time (RTT): The time between sending a packet and the arrival of the corresponding acknowledgement from the server

- Loss rate: The number of packets which did not reach the destination divided by the number of all sent packets

- ACK overhead: The additional number of the control messages sent by the TCP protocol compared to the number of message sent by the SCTP protocol.

2.3.4 Parameters

In the discussed simulation the following parameters have been used to show the different behaviour of the transport protocols:

- Type of transport protocol
  - Stream Control Transmission Protocol
    * Number of streams (1-5)
  - Transmission Control Protocol (Standard implementation of INET framework)
- Moving speed of mobile devices
  - Speed (5, 10, 20, 30, 70 km/h)

2.4 Results

2.4.1 Expected results

In this section a prediction about the results of the discussed simulation approach is made.

Throughput   The TCP should have a slightly higher throughput than SCTP if only one stream is used. Because the SCTP can not bring in his strength (multiple streams, message oriented). Increasing the number of streams while having a small loss rate (e.g. slow velocity) is not enough to beat the performance of TCP. But if there are losses (e.g. high velocity), SCTP has the mechanism to achieve a higher throughput.

Loss rate   A small moving speed of the mobile devices should result in a small loss rate. This behaviour is expected because the routing needs enough time to find a new route to the destination. If the velocity will be increased, then the loss rate increases too. The type of the transport protocol should have no influence to the loss rate.

Round trip time (RTT)   Using only one stream of the SCTP will not reduce the RTT and therefore the TCP will have a smaller one. But if multiple stream are used then the SCTP should reduce this metric significantly. The parallelism could be one reason to explain this. Generally, a higher velocity of the devices results in higher RTT (OLSR has to find new routes faster).

ACK Overhead   Increasing velocity will cause an increasing ACK overhead of TCP because packets are dropped and have to be retransmitted. The number of packet requests of SCTP increases with a higher number of streams and increasing velocity because of packet loss.

2.4.2 Measured results

Throughput   First of all the throughput of TCP and STCP have been compared with different movement velocities.
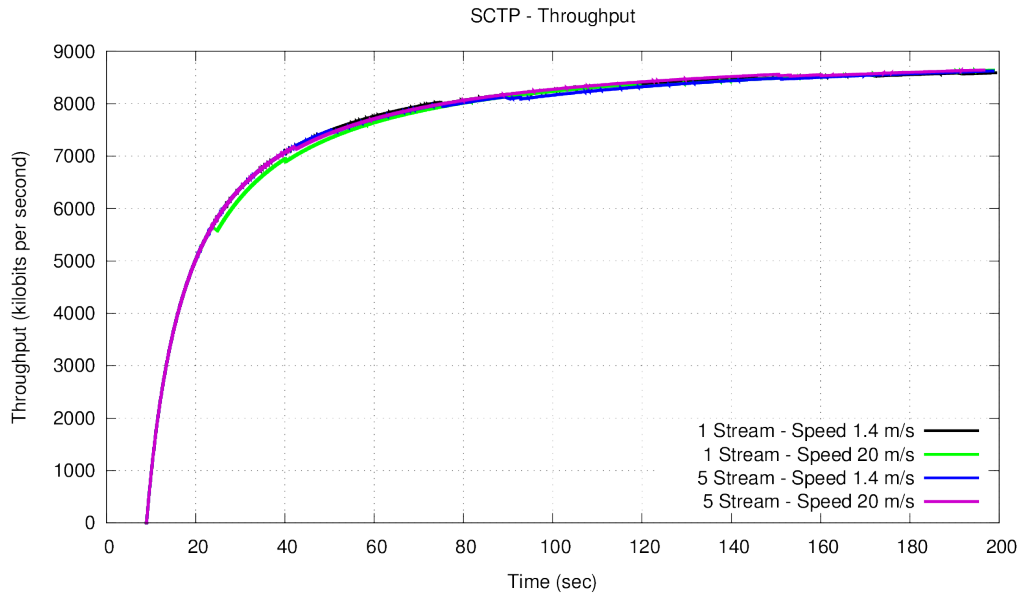
Figure 2: SCTP throughput under varying device movement

Figure 2 shows the throughput over time with 1 and 5 streams and the velocities 1.4 m/s and 20 m/s. Unfortunately the expected performance loss with a higher velocity is not be visible. Independent from the number of streams and the velocity the throughput over time is always in the same range. Because data is transferred it can be assumed that the ad hoc routing is working. Nevertheless it might be possible that the wireless range of the sender and the receiver is enough big that no more hop destinations between are necessary. To validate this assumption the throughput of TCP has been considered.
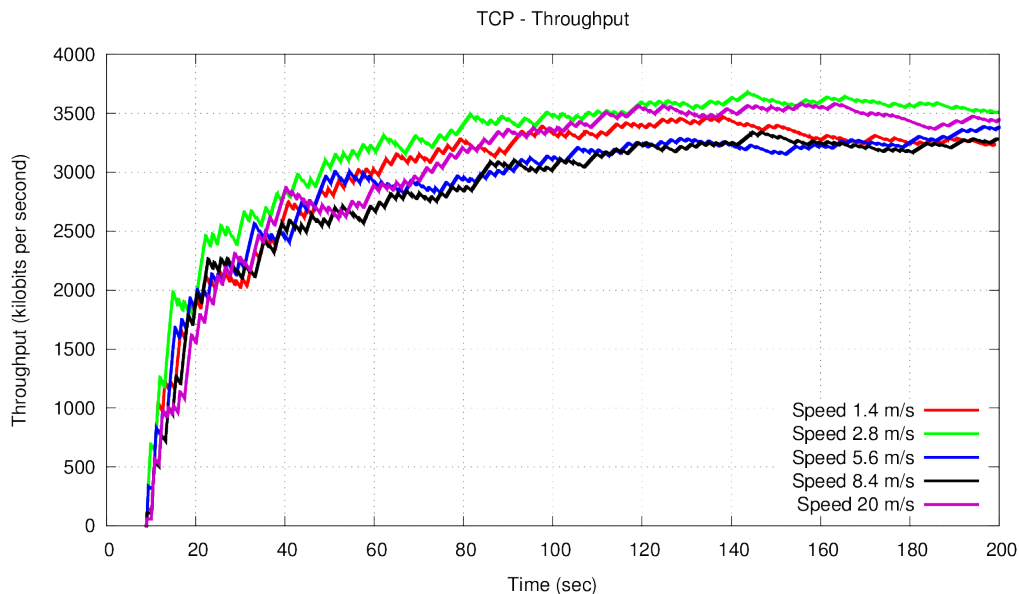

Figure 3: TCP throughput under varying device movement

Figure 3 compares the throughput of the TCP file transfer with different speeds. Like already seen in Figure 2 no direct influence of the increasing velocity can be seen. For example the results of velocity 2.8 and 20 m/s are very close and do not differ significantly like expected. On the basis of these both results the previ-

Network Performance and Modelling
FS2013

Practical Exercise
Final Report

Steve Aschwanden 05-480-686
Tobias Schmid 07-909-294

ous announced apprehension that hops between sender and receiver are not used for packet transmission has come true. Modifying the simulation parameters like increasing the path-loss parameter alpha or reducing the maximum transmission power lead to no result because packet transmission was not even possible.

To still be able to discuss the obtained result the velocity parameter will not considered in the following metrics and graphics.
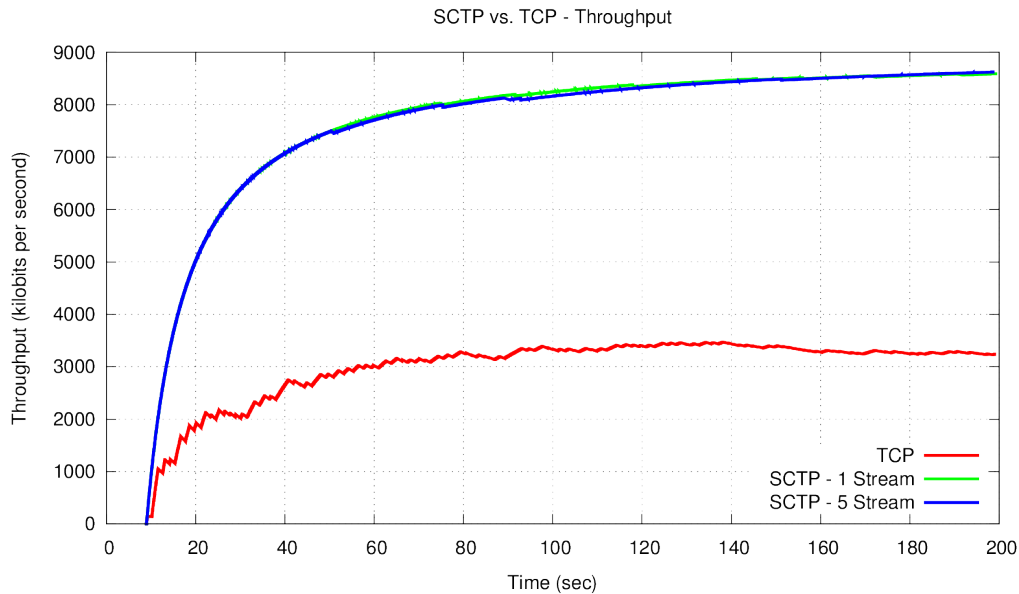

Figure 4: TCP vs. SCTP throughput

Figure 4 compares the throughput of TCP and SCTP over time. It can be seen that independent of the transport protocol the throughput starts growing after about 9 seconds of simulation time. Even thought the file transfer has been initiated at simulation start time this delay exists. A reason for that is the ad hoc routing using the OLSR protocol. While applying this protocol messages can only be forwarded when a corresponding route to the given destination has been found. OLSR needs some time to find neighbours and setup the routing tables. Finally when a route has been found data will be sent and the throughput increases with increasing time. In the case of SCTP the throughput increases logarithmic until about the 120 simulation second and then stays constant at the maximum value. In contrast to the smooth curve of SCTP, the throughput increase of TCP is rather erratic. This behaviour is caused by the flow and congestion control mechanisms of the TCP protocol. After about 80 simulation seconds the maximum throughput is reached and stays nearly constant until the end of the simulation.

Round trip time (RTT)   To compare the round trip time of both of these transport protocols the already in INET implemented statistical vector RTT of the associated protocol has been used. Figure 5 shows a comparison of the obtained results. Unfortunately the unit of the statistical vector is not documented so it was assumed that the values are store in seconds and have been downscaled to milliseconds later. The comparison shows that TCP has values with at most single digit millisecond range where SCTP with one stream has much higher values between 50 and 300ms. The values for SCTP with 5 streams have a strange behaviour the values are in most of the cases 1000ms with some statistical outliers.
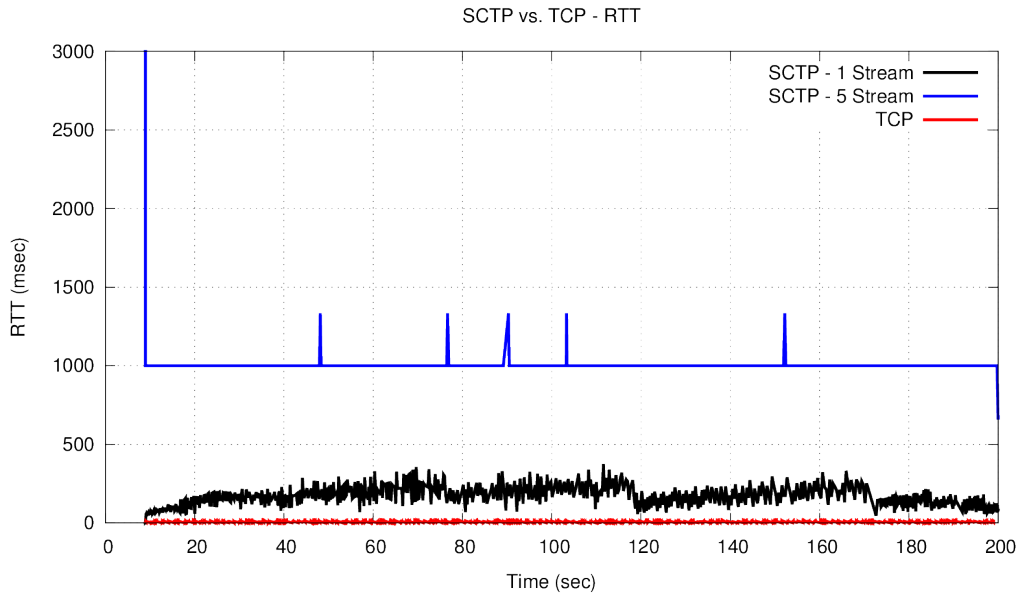
Figure 5: TCP vs. SCTP round trip time (RTT)

There are two possibilities to explain this behaviour:

1. The higher number of streams increases the occupancy of the wireless medium and generates a higher number of packet collisions which are handled on the MAC layer. For this reason requested packets will arrive later and the RTT will be higher.

2. The RTT metric implementation in the INET Framework is bogus.

Loss Rate   The loss rate of TCP has been calculated using the the number of sent seq packets on the server side and the number of received seq packets on the client side. Table 2.4.2 shows the the results of the three runs and the average of these results.

Speed	# Sent Packets	# Received Packets	Loss Rate
1.4 m/s	150788	150276	0.34 %
8.4 m/s	153413	152853	0.37 %
20 m/s	147834	147325	0.34 %
Average	-	-	0.35 %

Table 2: Loss rate TCP simulation runs

ACK Overhead

3   Conclusion

OMNeT++ is a powerful network simulation tool with many extensions for almost every technology and situation. But during the project the framework has shown some drawbacks. The behaviour of the lower layers is hard to understand and the insertion of new statistical values needs a lot of effort. The settings in the configuration file have a big influence to these. For example the creation time of a packet isn't available in some configurations and continuative calculations not possible. Another problem is the unavailable support of the community. There are some board entries, majoritarian without a solution. Therefore the user has to try different settings and spent a lot of time to solve the problem. But the included examples are very helpful.

The documentation of the INET extension is insufficient and the meanings of different parameters are unclear or not described. Only a look at the source code can help, but the whole function is mostly distributed over many files. The amount of data which will be created by the application is huge. The post-processing needs patience and the right tools. The export of the results is not very user friendly. Another missing part is a testing environment where the behaviour of newly generated modules can be tested or the behaviour of existing modules can be visualized.

References

[1] SCTP versus TCP: Comparing the Performance of Transport Protocols for Web Traffic, Rajesh Rajamani, Sumit Kumar, Nikhil Gupta, Computer Sciences Department, University of Wisconsin-Madison, July 2002, `http://pages.cs.wisc.edu/~sumit/extlinks/sctp.pdf`

[2] A comparison of TCP and SCTP performance using the HTTP protocol, Henrik Osterdahl, `http://www.ict.kth.se/courses/IK1550/Sample-papers/Henrik_Oesterdahl-sctp-20050525.pdf`

[3] OMNeT++ Network Simulation Framework, `http://www.omnetpp.org/`

[4] Introduction to the Stream Control Transmission Protocol (SCTP), Paul Stalvig, `http://www.f5.com/pdf/white-papers/sctp-introduction-wp.pdf`

[5] Comparing AODV and OLSR Routing Protocols, Aleksandr Huhtonen, Helsinki University of Technology, Telecommunication Software and Multimedia Laboratory, `http://edge.cs.drexel.edu/regli/Classes/CS680/Papers/Ad%20Hoc/Routing/aodv-oslr-comparison.pdf`